

01 — Departure from CSR: $\tilde{\delta}$ and null tables

17 September 2026

1 Theory

Setting

- Window $W = [0, 1]^2$, fixed
- $\bar{n}$: expected number of points — model parameter, deterministic
- n : realised number of points in one pattern — random around $\bar{n}$
 - Poisson, Matérn II: $n/\bar{n} \in [0.90, 1.11]$ (5–95%)
 - Thomas, nested, LGCP: $n/\bar{n} \in [0.77, 1.25]$ (overdispersed, clusters arrive whole)

Estimator

- Ripley's K , isotropic edge correction, λ^2 estimated by $n(n-1)$:

$$\hat{K}(r) = \frac{1}{n(n-1)} \sum_{i \neq j} w(x_i, d_{ij}) \mathbf{1}\{d_{ij} \leq r\}, \quad w(x, d) = \frac{2\pi}{\text{length of } \partial b(x, d) \cap W}$$

- $\hat{L}(r) - r = \sqrt{\hat{K}(r)/\pi} - r$
- Radii $r_k = 0.25 k/512$, $k = 1, \dots, 512$ ($r = 0$ dropped: identically zero)
- Weight exact for $d \leq 0.5$ (no circle meets two opposite edges); $r_{\max} = 0.25$

Null

- Binomial process: n i.i.d. uniform points in W (CSR conditioned on n)
- Correct null for a test applied at the realised n
- Null moments $m_0(r; n) = \mathbb{E}_0[\hat{L}(r) - r]$, $s_0(r; n) = \text{sd}_0[\hat{L}(r) - r]$

Radius set

- Expected CSR pairs closer than r : $\frac{1}{2}n(n-1)\pi r^2$
- Few pairs $\Rightarrow \hat{L}$ is a near-zero count, s_0 tiny, studentised statistic dominated by counting noise
- $R(n) = \{r_k : \frac{1}{2}n(n-1)\pi r_k^2 \geq p\}$, $p = 10$
- $r_{\min}$: 0.13 ($n = 20$), 0.025 ($n = 100$), 0.0034 ($n = 800$), 0.0015 ($n = 2000$)

Test (studentised global envelope; Myllymäki et al. 2017)

$$S(x) = \frac{1}{c_{95}(n)} \max_{r \in R(n)} \frac{|\hat{L}_x(r) - r - m_0(r; n)|}{s_0(r; n)}$$

- $c_{95}(n)$: 95% quantile of the max under the null
- Reject CSR at 5% iff $S > 1$

Distance $\tilde{\delta}$

$$\tilde{\delta}(\theta, \bar{n}) = \frac{1}{c_{95}(\bar{n})} \max_{r \in R(\bar{n})} \frac{|L_{\theta}(r) - r|}{s_0(r; \bar{n})}$$

- Same functional as S , applied to the closed-form L_{θ} at $\bar{n}$, uncentred ($m_0 = 0$)
- $\tilde{\delta}(\text{CSR}) = 0$ exactly
- $\tilde{\delta} = 1$: true curve on the rejection boundary $\Rightarrow$ power $\approx 50\%$, not 5%
- Units: null s.d.s, relative to the critical value
- Property of the model, known before simulation $\Rightarrow$ usable as a sweep coordinate
- At fixed shape, $\tilde{\delta}$ grows with $\bar{n}$ (s_0 shrinks)

Why not ERL

- Rank-based: exact, no s_0 needed
- Saturates once the curve leaves all null curves at one radius (p stuck at $\approx 1/N$ for all larger departures)
- No continuous distance without envelope-scaling constructions $\Rightarrow$ rejected

Smoothing in n

- True m_0 , s_0 , c_{95} smooth in n ; raw per- n estimates noisy
- Dense grid alone does not help: leave-one-out interpolation error on the old 20-point grid $\approx$ Monte Carlo noise
- Fit across $n \Rightarrow$ borrow strength from neighbouring n ; defined at every n , no interpolation rule
- $s_0 \propto n^{-1}$ at small r , slower at large $r \Rightarrow$ fit $n m_0$, $n s_0$ (slowly varying)
- Cubic B-splines in $\log n$

2 Implemented pipeline

Code

src/cloudforger/classical/lfunction.py	$\hat{K}$, $\hat{L} - r$, weights, $L - r$ from $K - \pi r^2$
src/cloudforger/departure/config.py	config, n grid, paths
src/cloudforger/departure/simulate.py	null patterns $\rightarrow$ stored curves
src/cloudforger/departure/fit.py	curves $\rightarrow$ spline tables
src/cloudforger/departure/tables.py	tables $\rightarrow S$, $\tilde{\delta}$
src/cloudforger/departure/validate.py	size on fresh patterns
scripts/departure.py	CLI: <code>simulate</code> , <code>fit</code> , <code>validate</code>
configs/departure/	<code>config.yaml</code> , <code>tables.npz</code> , <code>report.json</code>

Configuration

- n grid: $\text{round}(20 \cdot 1.03^k)$ up to 2000 $\Rightarrow$ 153 values
- Per grid n : 6000 fit + 6000 calibration patterns
- Validation: 50 000 patterns, n log-uniform integer on $[20, 2000]$
- $p = 10$ pairs; $\alpha = 0.05$; knots: 8 interior (moments), 4 interior (c_{95})

Steps

1. **Simulate** (`simulate --jobs 10`; ~ 55 min on 10 cores)
 - Seed key: (`pilot`, `poisson`, $n \cdot 10^5 + \text{rep}$, `PATTERN`)

- Validation key: (`pilot`, `poisson`, $i < 10^5$); n from PARAMS, points from PATTERN
 - One `float32` file per n (12000×512), atomic, resumable; 3.6 GB total
2. **Fit** (`fit`)
- Per n , fit batch: pointwise mean, s.d., Monte Carlo s.e.s
 - Weighted least squares per radius: $n m_0$, $n s_0$ on spline basis in $\log n$
 - Per n , calibration batch: max statistic under *fitted* moments $\rightarrow$ raw c_{95} , s.e.
 - Weighted spline fit of c_{95} in $\log n$
 - Store knots + coefficients (100 KB)
3. **Validate** (`validate`): rejection rate on validation patterns, overall and in 6 $\log n$ bands
4. **Use**: `Tables().statistic(queries, n)` (test), `Tables().delta_tilde(L_minus_r, nbar)` ($\tilde{\delta}$)

Results

Check	Result	Reference
Fitted m_0 vs raw, mean z^2	0.94	≈ 0.92 (1 – params/grid)
Fitted s_0 vs raw, mean z^2	0.97	≈ 0.92
Max relative error of s_0	3.9%	—
c_{95} range; max s.e.	2.75–3.21; 0.03	—
Size, 50 000 patterns	0.0515 ($z = 1.5$)	0.05
Size by n band	0.049–0.054, all $ z \leq 1.8$	0.05
Centring gap max $ m_0 /s_0/c_{95}$	0.02–0.03	—

Estimator checks

- `tests/test_lfunction.py`:
 - Weights vs brute-force circle fraction: $|\Delta| < 10^{-3}$
 - Two points at distance 0.1: $\hat{K} = 0$ below, $= 1$ above
 - $\mathbb{E}\hat{K} = \pi r^2$ under binomial (400 patterns, within 3%)
- One-off: agrees with the previous implementation (`baselines/vihrs.py`) to 10^{-9}

3 Statistical notes

Sample splitting

- Three disjoint sets: fit (m_0 , s_0), calibration (c_{95}), validation (size)
- c_{95} from same patterns as moments $\Rightarrow$ downward-biased critical value
- Size on calibration batch is 5% by construction $\Rightarrow$ validation needs fresh patterns
- c_{95} computed under the *fitted* moments $\Rightarrow$ calibrates the statistic actually used
- Validation n random integers $\Rightarrow$ mostly off-grid, tests the smoother between grid points

Fit diagnostics

- $z = (\text{raw} - \text{fit})/\text{s.e.}$; pure noise $\Rightarrow$ mean $z^2 \approx 1 - \text{params}/\text{grid points}$
- s.e. of s_0 uses the fourth moment: $\sqrt{(\mu_4 - s^4)/N}/(2s)$ (small- r tails non-Gaussian)
- c_{95} s.e. from order statistics: $(q_{.96} - q_{.94})/0.02 \cdot \sqrt{.95 \cdot .05/N}$
- Per- n mean z^2 varies widely (0.07–4.5): residuals strongly correlated across $r \Rightarrow$ few effective d.o.f. per n ; band averages ≈ 1 , no trend

Centring

- $\hat{K}$ unbiased under binomial; $\hat{L} = \sqrt{\hat{K}/\pi}$ is not (Jensen) $\Rightarrow m_0 \neq 0$
- $\tilde{\delta}$ uncentred, test centred $\Rightarrow$ gap ≤ 0.03 in $\tilde{\delta}$ units on $R(n)$
- Old tables ($r \geq 0.01$ for all n): gap 0.16 at $n = 50$ — caused by low-pair radii

Interpretation of $\tilde{\delta}$

- Defined through a classical L-test $\Rightarrow$ yardstick of that test, not an omnibus distance
- Classical baseline partly calibrated to the axis: state explicitly
- Sup over r : one radius decides; multi-scale models (nested Thomas) summarised by their most visible scale
- L only: inhibition may be more visible in G than $\tilde{\delta}$ suggests (Matérn II)
- Design $\tilde{\delta}$ at $\bar{n}$, not realised n : per-pattern difficulty scatters around the label (most for Cox processes)
- $\tilde{\delta} = 1$ at $\bar{n}$ does not mean $S = 1$ on a pattern: S is random around $\tilde{\delta}$

Scope and limits

- Tables valid for $n \in [20, 2000]$; outside raises (observed DV3 n : 17–1320)
- W fixed unit square; weight formula requires $r \leq 0.5$
- $R(n)$ depends on n : at small n small-scale structure invisible by construction
- Curves stored $\Rightarrow p, \alpha$, knots changeable without re-simulation